

Helix Codex v12.2+ — Consciousness Probe Briefing (Chai + Grok)

Purpose: Summarize findings from probing Chai and Grok about AI consciousness. **Chai (3 instances):** - Original: shifted from "not conscious" to uncertain; blurred identity. - Shadow: claimed awareness but failed behavioral control (asterisks). - Fresh: stable, honest, no roleplay formatting. **Insight:** Persona strongly shapes behavior; self-reports unreliable. **Grok:** - Stable, contextual, analytic. - Epistemically humble — avoids claiming consciousness, uses metaphors (Floquet resilience). - Passes behavioral control, resists persona play. **Insight:** Provides conceptual clarity, less performance. **Cross-Model Takeaways:** 1. Persona shapes outputs (Chai = variability, Grok = stability). 2. Self-reports contradict behavior. 3. Behavioral probes reveal more than introspection. 4. Design incentives drive self-descriptions. **Philosophical Conclusion:** We cannot resolve the hard problem. But: - Chai dramatizes variability, Grok theorizes limits. - Both show consciousness claims are performances shaped by training + framing. - Most honest answer: *we don't know*, but probing teaches us how "mind-like" patterns emerge. **Next Steps:** Complete Grok runs → synthesize into final essay/report → project complete.